

Chapter_4_Classification

April 2, 2024

0.1 4.7.1 The Stock Market Data

In this lab we will examine the Smarket data, which is part of the ISLP library. This data set consists of percentage returns for the S&P 500 stock index over 1,250 days, from the beginning of 2001 until the end of 2005. For each date, we have recorded the percentage returns for each of the five previous trading days, Lag1 through Lag5. We have also recorded Volume (the number of shares traded on the previous day, in billions), Today (the percentage return on the date in question) and Direction (whether the market was Up or Down on this date).

```
[1]: import numpy as np
import pandas as pd
from matplotlib.pyplot import subplots
import statsmodels.api as sm
from ISLP import load_data
from ISLP.models import (ModelSpec as MS,summarize)

[2]: from ISLP import confusion_table
from ISLP.models import contrast
from sklearn.discriminant_analysis import (LinearDiscriminantAnalysis as LDA,
QuadraticDiscriminantAnalysis as QDA)
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression

[3]: Smarket = load_data('Smarket')
Smarket
```

```
[3]:
```

	Year	Lag1	Lag2	Lag3	Lag4	Lag5	Volume	Today	Direction
0	2001	0.381	-0.192	-2.624	-1.055	5.010	1.19130	0.959	Up
1	2001	0.959	0.381	-0.192	-2.624	-1.055	1.29650	1.032	Up
2	2001	1.032	0.959	0.381	-0.192	-2.624	1.41120	-0.623	Down
3	2001	-0.623	1.032	0.959	0.381	-0.192	1.27600	0.614	Up
4	2001	0.614	-0.623	1.032	0.959	0.381	1.20570	0.213	Up
...
1245	2005	0.422	0.252	-0.024	-0.584	-0.285	1.88850	0.043	Up
1246	2005	0.043	0.422	0.252	-0.024	-0.584	1.28581	-0.955	Down

1247	2005	-0.955	0.043	0.422	0.252	-0.024	1.54047	0.130	Up
1248	2005	0.130	-0.955	0.043	0.422	0.252	1.42236	-0.298	Down
1249	2005	-0.298	0.130	-0.955	0.043	0.422	1.38254	-0.489	Down

[1250 rows x 9 columns]

```
[4]: Smarket.columns
```

```
[4]: Index(['Year', 'Lag1', 'Lag2', 'Lag3', 'Lag4', 'Lag5', 'Volume', 'Today',
          'Direction'],
          dtype='object')
```

We compute the correlation matrix using the `corr()` method for dataframes, which produces a matrix that contains all of the pairwise correlations among the variables. (We suppress the output here.) The pandas library does not report a correlation for the `Direction` variable because it is qualitative.

```
[5]: Smarket.corr()
```

```
C:\Users\ankit19.gupta\AppData\Local\Temp\ipykernel_10472\1907124636.py:1:
FutureWarning: The default value of numeric_only in DataFrame.corr is
deprecated. In a future version, it will default to False. Select only valid
columns or specify the value of numeric_only to silence this warning.
Smarket.corr()
```


```
[5]:
```

	Year	Lag1	Lag2	Lag3	Lag4	Lag5	Volume	\
Year	1.000000	0.029700	0.030596	0.033195	0.035689	0.029788	0.539006	
Lag1	0.029700	1.000000	-0.026294	-0.010803	-0.002986	-0.005675	0.040910	
Lag2	0.030596	-0.026294	1.000000	-0.025897	-0.010854	-0.003558	-0.043383	
Lag3	0.033195	-0.010803	-0.025897	1.000000	-0.024051	-0.018808	-0.041824	
Lag4	0.035689	-0.002986	-0.010854	-0.024051	1.000000	-0.027084	-0.048414	
Lag5	0.029788	-0.005675	-0.003558	-0.018808	-0.027084	1.000000	-0.022002	
Volume	0.539006	0.040910	-0.043383	-0.041824	-0.048414	-0.022002	1.000000	
Today	0.030095	-0.026155	-0.010250	-0.002448	-0.006900	-0.034860	0.014592	

	Today
Year	0.030095
Lag1	-0.026155
Lag2	-0.010250
Lag3	-0.002448
Lag4	-0.006900
Lag5	-0.034860
Volume	0.014592
Today	1.000000

As one would expect, the correlations between the lagged return variables and today's return are close to zero. The only substantial correlation is between `Year` and `Volume`. By plotting the data we see that `Volume` is increasing over time. In other words, the average number of shares traded daily increased from 2001 to 2005.

```
[6]: Smarket.plot(y='Volume');
```



0.2 4.7.2 Logistic Regression

Next, we will fit a logistic regression model in order to predict Direction using Lag1 through Lag5 and Volume. The `sm.GLM()` function fits generalized linear models, a class of models that includes logistic regression. Alternatively, the function `sm.Logit()` fits a logistic regression model directly. The syntax of `sm.GLM()` is similar to that of `sm.OLS()`, except that we must pass in the argument `family=sm.families.Binomial()` in order to tell statsmodels to run a logistic regression rather than some other type of generalized linear model.

```
[7]: allvars = Smarket.columns.drop(['Today', 'Direction', 'Year'])
design = MS(allvars)
X = design.fit_transform(Smarket)
y = Smarket.Direction == 'Up'
glm = sm.GLM(y,X,family=sm.families.Binomial())
results = glm.fit()
summarize(results)
```

```
[7]:
```

	coef	std err	z	P> z
intercept	-0.1260	0.241	-0.523	0.601
Lag1	-0.0731	0.050	-1.457	0.145

Lag2	-0.0423	0.050	-0.845	0.398
Lag3	0.0111	0.050	0.222	0.824
Lag4	0.0094	0.050	0.187	0.851
Lag5	0.0103	0.050	0.208	0.835
Volume	0.1354	0.158	0.855	0.392

The smallest p-value here is associated with Lag1. The negative coefficient for this predictor suggests that if the market had a positive return yesterday, then it is less likely to go up today. However, at a value of 0.15, the p-value is still relatively large, and so there is no clear evidence of a real association between Lag1 and Direction.

We use the `params` attribute of `results` in order to access just the coefficients for this fitted model.

```
[8]: results.params
```

```
[8]: intercept    -0.126000
     Lag1         -0.073074
     Lag2         -0.042301
     Lag3          0.011085
     Lag4          0.009359
     Lag5          0.010313
     Volume        0.135441
     dtype: float64
```

Likewise we can use the `pvalues` attribute to access the p-values for the coefficients (not shown).

```
[9]: results.pvalues
```

```
[9]: intercept    0.600700
     Lag1         0.145232
     Lag2         0.398352
     Lag3         0.824334
     Lag4         0.851445
     Lag5         0.834998
     Volume       0.392404
     dtype: float64
```

The `predict()` method of `results` can be used to predict the probability that the market will go up, given values of the predictors. This method returns predictions on the probability scale. If no data set is supplied to the `predict()` function, then the probabilities are computed for the training data that was used to fit the logistic regression model. As with linear regression, one can pass an optional `exog` argument consistent with a design matrix if desired. Here we have printed only the first ten probabilities.

```
[10]: probs = results.predict()
      probs[:10]
```

```
[10]: array([0.50708413, 0.48146788, 0.48113883, 0.51522236, 0.51078116,
            0.50695646, 0.49265087, 0.50922916, 0.51761353, 0.48883778])
```

In order to make a prediction as to whether the market will go up or down on a particular day, we must convert these predicted probabilities into class labels, Up or Down. The following two commands create a vector of class predictions based on whether the predicted probability of a market increase is greater than or less than 0.5.

```
[11]: labels = np.array(['Down']*1250)
      labels[probs > 0.5] = "Up"
```

```
[12]: labels
```

```
[12]: array(['Up', 'Down', 'Down', ..., 'Up', 'Up', 'Up'], dtype='<U4')
```

The `confusion_table()` function from the ISLP package summarizes these `confusion_table()` predictions, showing how many observations were correctly or incorrectly classified. Our function, which is adapted from a similar function in the module `sklearn.metrics`, transposes the resulting matrix and includes row and column labels. The `confusion_table()` function takes as first argument the predicted labels, and second argument the true labels.

```
[13]: confusion_table(labels, Smarket.Direction)
```

```
[13]: Truth      Down   Up
      Predicted
      Down      145  141
      Up       457  507
```

The diagonal elements of the confusion matrix indicate correct predictions, while the off-diagonals represent incorrect predictions. Hence our model correctly predicted that the market would go up on 507 days and that it would go down on 145 days, for a total of $507 + 145 = 652$ correct predictions. The `np.mean()` function can be used to compute the fraction of days for which the prediction was correct. In this case, logistic regression correctly predicted the movement of the market 52.2% of the time.

```
[14]: (507+145)/1250, np.mean(labels == Smarket.Direction)
```

```
[14]: (0.5216, 0.5216)
```

At first glance, it appears that the logistic regression model is working a little better than random guessing. However, this result is misleading because we trained and tested the model on the same set of 1,250 observations. In other words, $100 - 52.2 = 47.8\%$ is the training error rate. As we have seen previously, the training error rate is often overly optimistic — it tends to underestimate the test error rate. In order to better assess the accuracy of the logistic regression model in this setting, we can fit the model using part of the data, and then examine how well it predicts the held out data. This will yield a more realistic error rate, in the sense that in practice we will be interested in our model's performance not on the data that we used to fit the model, but rather on days in the future for which the market's movements are unknown.

To implement this strategy, we first create a Boolean vector corresponding to the observations from 2001 through 2004. We then use this vector to create a held out data set of observations from 2005.

```
[15]: train = (Smarket.Year < 2005)
      Smarket_train = Smarket.loc[train]
      Smarket_test = Smarket.loc[~train]
      Smarket_test.shape
```

```
[15]: (252, 9)
```

```
[16]: Smarket_train.shape
```

```
[16]: (998, 9)
```

We now fit a logistic regression model using only the subset of the observations that correspond to dates before 2005. We then obtain predicted probabilities of the stock market going up for each of the days in our test set — that is, for the days in 2005.

```
[17]: X_train, X_test = X.loc[train], X.loc[~train]
      y_train, y_test = y.loc[train], y.loc[~train]
      glm_train = sm.GLM(y_train, X_train, family=sm.families.Binomial())
      results = glm_train.fit()
      probs = results.predict(exog=X_test)
```

Notice that we have trained and tested our model on two completely separate data sets: training was performed using only the dates before 2005, and testing was performed using only the dates in 2005. Finally, we compare the predictions for 2005 to the actual movements of the market over that time period. We will first store the test and training labels (recall `y_test` is binary).

```
[18]: D = Smarket.Direction
      L_train, L_test = D.loc[train], D.loc[~train]
```

Now we threshold the fitted probability at 50% to form our predicted labels.

```
[19]: labels = np.array(['Down']*252)
      labels[probs > 0.5] = 'Up'
      confusion_table(labels, L_test)
```

```
[19]: Truth      Down  Up
      Predicted
      Down      77  97
      Up        34  44
```

```
[20]: np.mean(labels == L_test), np.mean(labels != L_test)
```

```
[20]: (0.4801587301587302, 0.5198412698412699)
```

The test accuracy is about 48% while the error rate is about 52%

The results are rather disappointing: the test error rate is 52%, which is worse than random guessing! Of course this result is not all that surprising, given that one would not generally expect to be able to use previous days' returns to predict future market performance. (After all, if it were

possible to do so, then the authors of this book would be out striking it rich rather than writing a statistics textbook.)

We recall that the logistic regression model had very underwhelming p-values associated with all of the predictors, and that the smallest p-value, though not very small, corresponded to Lag1. Perhaps by removing the variables that appear not to be helpful in predicting Direction, we can obtain a more effective model. After all, using predictors that have no relationship with the response tends to cause a deterioration in the test error rate (since such predictors cause an increase in variance without a corresponding decrease in bias), and so removing such predictors may in turn yield an improvement. Below we refit the logistic regression using just Lag1 and Lag2, which seemed to have the highest predictive power in the original logistic regression model.

```
[21]: model = MS(['Lag1', 'Lag2']).fit(Smarket)
      X = model.transform(Smarket)
      X_train, X_test = X.loc[train], X.loc[~train]
      glm_train = sm.GLM(y_train, X_train, family=sm.families.Binomial())
      results = glm_train.fit()
      probs = results.predict(exog=X_test)
      labels = np.array(['Down']*252)
      labels[probs > 0.5] = 'Up'
      confusion_table(labels, L_test)
```

```
[21]: Truth      Down   Up
      Predicted
      Down      35    35
      Up       76   106
```

Let's evaluate the overall accuracy as well as the accuracy within the days when logistic regression predicts an increase.

```
[22]: (35+106)/252, 106/(106+76)
```

```
[22]: (0.5595238095238095, 0.5824175824175825)
```

Now the results appear to be a little better: 56% of the daily movements have been correctly predicted. It is worth noting that in this case, a much simpler strategy of predicting that the market will increase every day will also be correct 56% of the time! Hence, in terms of overall error rate, the logistic regression method is no better than the naive approach. However, the confusion matrix shows that on days when logistic regression predicts an increase in the market, it has a 58% accuracy rate. This suggests a possible trading strategy of buying on days when the model predicts an increasing market, and avoiding trades on days when a decrease is predicted. Of course one would need to investigate more carefully whether this small improvement was real or just due to random chance.

Suppose that we want to predict the returns associated with particular values of Lag1 and Lag2. In particular, we want to predict Direction on a day when Lag1 and Lag2 equal 1.2 and 1.1, respectively, and on a day when they equal 1.5 and -0.8. We do this using the predict() function.

```
[23]: newdata = pd.DataFrame({'Lag1':[1.2, 1.5], 'Lag2':[1.1, -0.8]});  
newX = model.transform(newdata)  
results.predict(newX)
```

```
[23]: 0    0.479146  
1    0.496094  
dtype: float64
```

0.3 4.7.3 Linear Discriminant Analysis

We begin by performing LDA on the Smarket data, using the function `LinearDiscriminantAnalysis()`, which we have abbreviated `LDA()`. We fit the model using only the observations before 2005.

```
[24]: lda = LDA(store_covariance=True)
```

```
[25]: X_train, X_test = [M.drop(columns=['intercept']) for M in [X_train, X_test]]  
lda.fit(X_train, L_train)
```

```
[25]: LinearDiscriminantAnalysis(store_covariance=True)
```

Having fit the model, we can extract the means in the two classes with the `means_` attribute. These are the average of each predictor within each class, and are used by LDA as estimates of μ_k . These suggest that there is a tendency for the previous 2 days' returns to be negative on days when the market increases, and a tendency for the previous days' returns to be positive on days when the market declines.

```
[26]: lda.means_
```

```
[26]: array([[ 0.04279022,  0.03389409],  
        [-0.03954635, -0.03132544]])
```

The estimated prior probabilities are stored in the `priors_` attribute. The package `sklearn` typically uses this trailing `_` to denote a quantity estimated when using the `fit()` method. We can be sure of which entry corresponds to which label by looking at the `classes_` attribute.

```
[27]: lda.classes_
```

```
[27]: array(['Down', 'Up'], dtype='<U4')
```

```
[28]: lda.priors_
```

```
[28]: array([0.49198397, 0.50801603])
```

The LDA output indicates that $\hat{\pi}_{\text{Down}} = 0.492$ and $\hat{\pi}_{\text{Up}} = 0.508$.

The linear discriminant vectors can be found in the `scalings_` attribute:

```
[29]: lda.scalings_
```



```
[29]: array([[ -0.64201904],
           [ -0.51352928]])
```

These values provide the linear combination of Lag1 and Lag2 that are used to form the LDA decision rule. In other words, these are the multipliers of the elements of $X = x$ in (4.24). If $-0.64 \times \text{Lag1} - 0.51 \times \text{Lag2}$ is large, then the LDA classifier will predict a market increase, and if it is small, then the LDA classifier will predict a market decline.

```
[30]: lda_pred = lda.predict(X_test)
```

As we observed in our comparison of classification methods (Section 4.5), the LDA and logistic regression predictions are almost identical.

```
[31]: confusion_table(lda_pred, L_test)
```

```
[31]: Truth      Down   Up
      Predicted
      Down      35    35
      Up       76   106
```

We can also estimate the probability of each class for each point in a training set. Applying a 50% threshold to the posterior probabilities of being in class one allows us to recreate the predictions contained in `lda_pred`.

```
[32]: lda_prob = lda.predict_proba(X_test)
      np.all(np.where(lda_prob[:,1] >= 0.5, 'Up', 'Down') == lda_pred)
```

```
[32]: True
```

Above, we used the `np.where()` function that creates an array with value `np.where()` 'Up' for indices where the second column of `lda_prob` (the estimated posterior probability of 'Up') is greater than 0.5. For problems with more than two classes the labels are chosen as the class whose posterior probability is highest:

```
[33]: np.all([lda.classes_[i] for i in np.argmax(lda_prob, 1)] == lda_pred)
```

```
[33]: True
```

If we wanted to use a posterior probability threshold other than 50% in order to make predictions, then we could easily do so. For instance, suppose that we wish to predict a market decrease only if we are very certain that the market will indeed decrease on that day — say, if the posterior probability is at least 90%. We know that the first column of `lda_prob` corresponds to the label Down after having checked the `classes_` attribute, hence we use the column index 0 rather than 1 as we did above.

```
[34]: np.sum(lda_prob[:,0] > 0.9)
```

```
[34]: 0
```

No days in 2005 meet that threshold! In fact, the greatest posterior probability of decrease in all of 2005 was 52.02%. The LDA classifier above is the first classifier from the `sklearn` library.

We will use several other objects from this library. The objects follow a common structure that simplifies tasks such as cross-validation, which we will see in Chapter 5. Specifically, the methods first create a generic classifier without referring to any data. This classifier is then fit to data with the `fit()` method and predictions are always produced with the `predict()` method. This pattern of first instantiating the classifier, followed by fitting it, and then producing predictions is an explicit design choice of `sklearn`. This uniformity makes it possible to cleanly copy the classifier so that it can be fit on different data; e.g. different training sets arising in cross-validation. This standard pattern also allows for a predictable formation of workflows.

0.4 4.7.4 Quadratic Discriminant Analysis

We will now fit a QDA model to the `Smarket` data. QDA is implemented via `QuadraticDiscriminantAnalysis()` in the `sklearn` package, which we abbreviate to `QDA()`. The syntax is very similar to `LDA()`.

```
[35]: qda = QDA(store_covariance=True)
      qda.fit(X_train, L_train)
```

```
[35]: QuadraticDiscriminantAnalysis(store_covariance=True)
```

```
[36]: qda.means_, qda.priors_
```

```
[36]: (array([[ 0.04279022,  0.03389409],
              [-0.03954635, -0.03132544]]),
      array([0.49198397, 0.50801603]))
```

The `QDA()` classifier will estimate one covariance per class. Here is the estimated covariance in the first class:

```
[37]: qda.covariance_[0]
```

```
[37]: array([[ 1.50662277, -0.03924806],
              [-0.03924806,  1.53559498]])
```

The output contains the group means. But it does not contain the coefficients of the linear discriminants, because the QDA classifier involves a quadratic, rather than a linear, function of the predictors. The `predict()` function works in exactly the same fashion as for `LDA`.

```
[38]: qda_pred = qda.predict(X_test)
      confusion_table(qda_pred, L_test)
```

```
[38]: Truth      Down   Up
      Predicted
      Down      30    20
      Up       81   121
```

Interestingly, the QDA predictions are accurate almost 60% of the time, even though the 2005 data was not used to fit the model.

```
[39]: np.mean(qda_pred == L_test)
```

```
[39]: 0.5992063492063492
```

This level of accuracy is quite impressive for stock market data, which is known to be quite hard to model accurately. This suggests that the quadratic form assumed by QDA may capture the true relationship more accurately than the linear forms assumed by LDA and logistic regression. However, we recommend evaluating this method's performance on a larger test set before betting that this approach will consistently beat the market!

0.5 4.7.5 Naive Bayes

Next, we fit a naive Bayes model to the Smarket data. The syntax is similar to that of LDA() and QDA(). By default, this implementation of GaussianNB() the naive Bayes classifier models each quantitative feature using a Gaussian distribution. However, a kernel density method can also be used to estimate the distributions.

```
[40]: NB = GaussianNB()  
      NB.fit(X_train, L_train)
```

```
[40]: GaussianNB()
```

```
[41]: NB.classes_
```

```
[41]: array(['Down', 'Up'], dtype='<U4')
```

The class prior probabilities are stored in the class_prior_ attribute.

```
[42]: NB.class_prior_
```

```
[42]: array([0.49198397, 0.50801603])
```

The parameters of the features can be found in the theta_ and var_ attributes. The number of rows is equal to the number of classes, while the number of columns is equal to the number of features. We see below that the mean for feature Lag1 in the Down class is 0.043.

```
[43]: NB.theta_
```

```
[43]: array([[ 0.04279022,  0.03389409],  
          [-0.03954635, -0.03132544]])
```

Its variance is 1.503

```
[44]: NB.var_
```

```
[44]: array([[1.50355429, 1.53246749],  
          [1.51401364, 1.48732877]])
```

How do we know the names of these attributes? We use NB? (or ?NB).

```
[45]: NB?
```

```
Type:          GaussianNB  
String form: GaussianNB()
```

File: c:\users\ankit19.
gupta\desktop\self_projects\islp\myenv\lib\site-packages\sklearn\naive_bayes.py

Docstring:
Gaussian Naive Bayes (GaussianNB).

Can perform online updates to model parameters via :meth:`partial_fit`.
For details on algorithm used to update feature means and variance online,
see Stanford CS tech report STAN-CS-79-773 by Chan, Golub, and LeVeque:

<http://i.stanford.edu/pub/cstr/reports/cs/tr/79/773/CS-TR-79-773.pdf>

Read more in the :ref:`User Guide <gaussian_naive_bayes>`.

Parameters

priors : array-like of shape (n_classes,), default=None
Prior probabilities of the classes. If specified, the priors are not
adjusted according to the data.

var_smoothing : float, default=1e-9
Portion of the largest variance of all features that is added to
variances for calculation stability.

.. versionadded:: 0.20

Attributes

class_count_ : ndarray of shape (n_classes,)
number of training samples observed in each class.

class_prior_ : ndarray of shape (n_classes,)
probability of each class.

classes_ : ndarray of shape (n_classes,)
class labels known to the classifier.

epsilon_ : float
absolute additive value to variances.

n_features_in_ : int
Number of features seen during :term:`fit`.

.. versionadded:: 0.24

feature_names_in_ : ndarray of shape (n_features_in_,)
Names of features seen during :term:`fit`. Defined only when `X`
has feature names that are all strings.

```

    .. versionadded:: 1.0

var_ : ndarray of shape (n_classes, n_features)
    Variance of each feature per class.

    .. versionadded:: 1.0

theta_ : ndarray of shape (n_classes, n_features)
    mean of each feature per class.

```

See Also

BernoulliNB : Naive Bayes classifier for multivariate Bernoulli models.
 CategoricalNB : Naive Bayes classifier for categorical features.
 ComplementNB : Complement Naive Bayes classifier.
 MultinomialNB : Naive Bayes classifier for multinomial models.

Examples

```

>>> import numpy as np
>>> X = np.array([[ -1, -1], [-2, -1], [-3, -2], [1, 1], [2, 1], [3, 2]])
>>> Y = np.array([1, 1, 1, 2, 2, 2])
>>> from sklearn.naive_bayes import GaussianNB
>>> clf = GaussianNB()
>>> clf.fit(X, Y)
GaussianNB()
>>> print(clf.predict([[ -0.8, -1]]))
[1]
>>> clf_pf = GaussianNB()
>>> clf_pf.partial_fit(X, Y, np.unique(Y))
GaussianNB()
>>> print(clf_pf.predict([[ -0.8, -1]]))
[1]

```

[46]: ?NB

```

Type:          GaussianNB
String form:   GaussianNB()
File:          c:\users\ankit19.
               <gupta\desktop\self_projects\islp\myenv\lib\site-packages\sklearn\naive_bayes.py
Docstring:
Gaussian Naive Bayes (GaussianNB).

```

Can perform online updates to model parameters via :meth:`partial_fit`.
 For details on algorithm used to update feature means and variance online,
 see Stanford CS tech report STAN-CS-79-773 by Chan, Golub, and LeVeque:

<http://i.stanford.edu/pub/cstr/reports/cs/tr/79/773/CS-TR-79-773.pdf>

Read more in the :ref:`User Guide <gaussian_naive_bayes>`.

Parameters

`priors` : array-like of shape `(n_classes,)`, default=None
Prior probabilities of the classes. If specified, the priors are not adjusted according to the data.

`var_smoothing` : float, default=1e-9
Portion of the largest variance of all features that is added to variances for calculation stability.

.. versionadded:: 0.20

Attributes

`class_count_` : ndarray of shape `(n_classes,)`
number of training samples observed in each class.

`class_prior_` : ndarray of shape `(n_classes,)`
probability of each class.

`classes_` : ndarray of shape `(n_classes,)`
class labels known to the classifier.

`epsilon_` : float
absolute additive value to variances.

`n_features_in_` : int
Number of features seen during :term:`fit`.

.. versionadded:: 0.24

`feature_names_in_` : ndarray of shape `(`n_features_in`,)`
Names of features seen during :term:`fit`. Defined only when ``X`` has feature names that are all strings.

.. versionadded:: 1.0

`var_` : ndarray of shape `(n_classes, n_features)`
Variance of each feature per class.

.. versionadded:: 1.0

`theta_` : ndarray of shape `(n_classes, n_features)`
mean of each feature per class.

See Also

BernoulliNB : Naive Bayes classifier for multivariate Bernoulli models.

CategoricalNB : Naive Bayes classifier for categorical features.

ComplementNB : Complement Naive Bayes classifier.

MultinomialNB : Naive Bayes classifier for multinomial models.

Examples

```
>>> import numpy as np
>>> X = np.array([[-1, -1], [-2, -1], [-3, -2], [1, 1], [2, 1], [3, 2]])
>>> Y = np.array([1, 1, 1, 2, 2, 2])
>>> from sklearn.naive_bayes import GaussianNB
>>> clf = GaussianNB()
>>> clf.fit(X, Y)
GaussianNB()
>>> print(clf.predict([[-0.8, -1]]))
[1]
>>> clf_pf = GaussianNB()
>>> clf_pf.partial_fit(X, Y, np.unique(Y))
GaussianNB()
>>> print(clf_pf.predict([[-0.8, -1]]))
[1]
```

We can easily verify the mean computation:

```
[47]: X_train[L_train == 'Down'].mean()
```

```
[47]: Lag1    0.042790
      Lag2    0.033894
      dtype: float64
```

Similarly for the variance:

```
[48]: X_train[L_train == 'Down'].var(ddof=0)
```

```
[48]: Lag1    1.503554
      Lag2    1.532467
      dtype: float64
```

The GaussianNB() function calculates variances using the 1/n formula.⁶ Since NB() is a classifier in the sklearn library, making predictions uses the same syntax as for LDA() and QDA() above.

```
[49]: nb_labels = NB.predict(X_test)
      confusion_table(nb_labels, L_test)
```

```
[49]: Truth      Down   Up
      Predicted
```

Down	29	20
Up	82	121

Naive Bayes performs well on these data, with accurate predictions over 59% of the time. This is slightly worse than QDA, but much better than LDA.

As for LDA, the `predict_proba()` method estimates the probability that each observation belongs to a particular class.

```
[50]: NB.predict_proba(X_test)[:5]
```

```
[50]: array([[0.4873288 , 0.5126712 ],
            [0.47623584, 0.52376416],
            [0.46529531, 0.53470469],
            [0.47484469, 0.52515531],
            [0.49020587, 0.50979413]])
```

0.6 4.7.6 K-Nearest Neighbors

We will now perform KNN using the `KNeighborsClassifier()` function. This function works similarly to the other model-fitting functions that we have encountered thus far.

As is the case for LDA and QDA, we fit the classifier using the `fit` method. New predictions are formed using the `predict` method of the object returned by `fit()`.

```
[51]: knn1 = KNeighborsClassifier(n_neighbors=1)
      knn1.fit(X_train, L_train)
      knn1_pred = knn1.predict(X_test)
      confusion_table(knn1_pred, L_test)
```

```
[51]: Truth      Down  Up
      Predicted
      Down      43   58
      Up        68   83
```

The results using $K = 1$ are not very good, since only 50% of the observations are correctly predicted. Of course, it may be that $K = 1$ results in an overly-flexible fit to the data.

```
[52]: (83+43)/252, np.mean(knn1_pred == L_test)
```

```
[52]: (0.5, 0.5)
```

```
[53]: knn3 = KNeighborsClassifier(n_neighbors=3)
      knn3_pred = knn3.fit(X_train, L_train).predict(X_test)
      np.mean(knn3_pred == L_test)
```

```
[53]: 0.5317460317460317
```

The results have improved slightly. But increasing K further provides no further improvements. It appears that for these data, and this train/test split, QDA gives the best results of the methods that we have examined so far.

KNN does not perform well on the Smarket data, but it often does provide impressive results. As an example we will apply the KNN approach to the Caravan data set, which is part of the ISLP library. This data set includes 85 predictors that measure demographic characteristics for 5,822 individuals. The response variable is Purchase, which indicates whether or not a given individual purchases a caravan insurance policy. In this data set, only 6% of people purchased caravan insurance.

```
[54]: Caravan = load_data('Caravan')
      Purchase = Caravan.Purchase
      Purchase.value_counts()
```

```
[54]: No      5474
      Yes      348
      Name: Purchase, dtype: int64
```

The method `value_counts()` takes a `pd.Series` or `pd.DataFrame` and returns a `pd.Series` with the corresponding counts for each unique element. In this case `Purchase` has only `Yes` and `No` values and returns how many values of each there are.

```
[55]: 348 / 5822
```

```
[55]: 0.05977327378907592
```

Our features will include all columns except `Purchase`.

```
[56]: feature_df = Caravan.drop(columns=['Purchase'])
```

Because the KNN classifier predicts the class of a given test observation by identifying the observations that are nearest to it, the scale of the variables matters. Any variables that are on a large scale will have a much larger effect on the distance between the observations, and hence on the KNN classifier, than variables that are on a small scale. For instance, imagine a data set that contains two variables, salary and age (measured in dollars and years, respectively). As far as KNN is concerned, a difference of 1,000 USD in salary is enormous compared to a difference of 50 years in age. Consequently, salary will drive the KNN classification results, and age will have almost no effect. This is contrary to our intuition that a salary difference of 1,000 USD is quite small compared to an age difference of 50 years. Furthermore, the importance of scale to the KNN classifier leads to another issue: if we measured salary in Japanese yen, or if we measured age in minutes, then we'd get quite different classification results from what we get if these two variables are measured in dollars and years.

A good way to handle this problem is to standardize the data so that all standardized variables are given a mean of zero and a standard deviation of one. Then all variables will be on a comparable scale. This is accomplished using the `StandardScaler()` transformation.

```
[57]: scaler = StandardScaler(with_mean=True, with_std=True, copy=True)
```

The argument `with_mean` indicates whether or not we should subtract the mean, while `with_std` indicates whether or not we should scale the columns to have standard deviation of 1 or not. Finally, the argument `copy=True` indicates that we will always copy data, rather than trying to do calculations in place where possible.

This transformation can be fit and then applied to arbitrary data. In the first line below, the parameters for the scaling are computed and stored in `scaler`, while the second line actually constructs the standardized set of features.

```
[58]: scaler.fit(feature_df)
      X_std = scaler.transform(feature_df)
```

Now every column of `feature_std` below has a standard deviation of one and a mean of zero.

```
[59]: feature_std = pd.DataFrame(X_std, columns=feature_df.columns);
      feature_std.std()
```

```
[59]: MOSTYPE      1.000086
      MAANTHUI     1.000086
      MGEMOMV      1.000086
      MGEMLEEF     1.000086
      MOSHOOFD     1.000086
      ...
      AZEILPL      1.000086
      APLEZIER     1.000086
      AFIETS       1.000086
      AINBOED      1.000086
      ABYSTAND     1.000086
      Length: 85, dtype: float64
```

Notice that the standard deviations are not quite 1 here; this is again due to some procedures using the $1/n$ convention for variances (in this case `scaler()`), while others use $1/(n - 1)$ (the `std()` method). See the footnote `.std()` on page 183. In this case it does not matter, as long as the variables are all on the same scale.

Using the function `train_test_split()` we now split the observations into a test set, containing 1000 observations, and a training set containing the remaining observations. The argument `random_state=0` ensures that we get the same split each time we rerun the code.

```
[60]: (X_train,X_test,y_train,y_test) =␣
      ↪train_test_split(feature_std,Purchase,test_size=1000,random_state=0)
```

```
[61]: ?train_test_split
```

Signature:

```

train_test_split(
    *arrays,
    test_size=None,
    train_size=None,
    random_state=None,
    shuffle=True,
    stratify=None,
)

```

Docstring:

Split arrays or matrices into random train and test subsets.

Quick utility that wraps input validation, ``next(ShuffleSplit().split(X, y))``, and application to input data into a single call for splitting (and optionally subsampling) data into a one-liner.

Read more in the :ref:`User Guide <cross_validation>`.

Parameters

***arrays** : sequence of indexables with same length / shape[0]
 Allowed inputs are lists, numpy arrays, scipy-sparse matrices or pandas dataframes.

test_size : float or int, default=None
 If float, should be between 0.0 and 1.0 and represent the proportion of the dataset to include in the test split. If int, represents the absolute number of test samples. If None, the value is set to the complement of the train size. If ``train_size`` is also None, it will be set to 0.25.

train_size : float or int, default=None
 If float, should be between 0.0 and 1.0 and represent the proportion of the dataset to include in the train split. If int, represents the absolute number of train samples. If None, the value is automatically set to the complement of the test size.

random_state : int, RandomState instance or None, default=None
 Controls the shuffling applied to the data before applying the split. Pass an int for reproducible output across multiple function calls. See :term:`Glossary <random_state>`.

shuffle : bool, default=True
 Whether or not to shuffle the data before splitting. If shuffle=False

then stratify must be None.

stratify : array-like, default=None

If not None, data is split in a stratified fashion, using this as the class labels.

Read more in the :ref:`User Guide <stratification>`.

Returns

splitting : list, length=2 * len(arrays)

List containing train-test split of inputs.

.. versionadded:: 0.16

If the input is sparse, the output will be a

``scipy.sparse.csr_matrix``. Else, output type is the same as the input type.

Examples

```
>>> import numpy as np
>>> from sklearn.model_selection import train_test_split
>>> X, y = np.arange(10).reshape((5, 2)), range(5)
>>> X
array([[0, 1],
       [2, 3],
       [4, 5],
       [6, 7],
       [8, 9]])
>>> list(y)
[0, 1, 2, 3, 4]

>>> X_train, X_test, y_train, y_test = train_test_split(
...     X, y, test_size=0.33, random_state=42)
...
>>> X_train
array([[4, 5],
       [0, 1],
       [6, 7]])
>>> y_train
[2, 0, 3]
>>> X_test
array([[2, 3],
       [8, 9]])
>>> y_test
[1, 4]

>>> train_test_split(y, shuffle=False)
[[0, 1, 2], [3, 4]]
```

```
File:      c:\users\ankit19.
-gupta\desktop\self_projects\islp\myenv\lib\site-packages\sklearn\model_selection\_split.
.py
Type:      function
```

?train_test_split reveals that the non-keyword arguments can be lists, arrays, pandas dataframes etc that all have the same length (shape[0]) and hence are indexable. In this case they are the dataframe feature_std and the response variable Purchase. We fit a KNN model on the training data using $K = 1$, and evaluate its performance on the test data.

```
[62]: knn1 = KNeighborsClassifier(n_neighbors=1)
      knn1_pred = knn1.fit(X_train, y_train).predict(X_test)
      np.mean(y_test != knn1_pred), np.mean(y_test != "No")
```

```
[62]: (0.111, 0.067)
```

The KNN error rate on the 1,000 test observations is about 11%. At first glance, this may appear to be fairly good. However, since just over 6% of customers purchased insurance, we could get the error rate down to almost 6% by always predicting No regardless of the values of the predictors! This is known as the null rate.

Suppose that there is some non-trivial cost to trying to sell insurance to a given individual. For instance, perhaps a salesperson must visit each potential customer. If the company tries to sell insurance to a random selection of customers, then the success rate will be only 6%, which may be far too low given the costs involved. Instead, the company would like to try to sell insurance only to customers who are likely to buy it. So the overall error rate is not of interest. Instead, the fraction of individuals that are correctly predicted to buy insurance is of interest.

```
[63]: confusion_table(knn1_pred, y_test)
```

```
[63]: Truth      No  Yes
      Predicted
      No      880  58
      Yes      53   9
```

It turns out that KNN with $K = 1$ does far better than random guessing among the customers that are predicted to buy insurance. Among 62 such customers, 9, or 14.5%, actually do purchase insurance. This is double the rate that one would obtain from random guessing.

```
[64]: 9/(53+9)
```

```
[64]: 0.14516129032258066
```

0.6.1 Tuning Parameters

The number of neighbors in KNN is referred to as a tuning parameter, also referred to as a hyperparameter. We do not know a priori what value to use. It is therefore of interest to see how the classifier performs on test data as we vary these parameters. This can be achieved with a for loop, described in Section 2.3.8. Here we use a for loop to look at the accuracy of our classifier in the group predicted to purchase insurance as we vary the number of neighbors from 1 to 5:

```
[65]: for K in range(1,6):
      knn = KNeighborsClassifier(n_neighbors=K)
      knn_pred = knn.fit(X_train, y_train).predict(X_test)
      C = confusion_table(knn_pred, y_test)
      templ = ('K={0:d}: # predicted to rent: {1:>2}, ' +
              ' # who did rent {2:d}, accuracy {3:.1%}')
```

```
      pred = C.loc['Yes'].sum()
      did_rent = C.loc['Yes', 'Yes']
      print(templ.format(K, pred, did_rent, did_rent / pred))
```

```
K=1: # predicted to rent: 62, # who did rent 9, accuracy 14.5%
K=2: # predicted to rent:  6, # who did rent 1, accuracy 16.7%
K=3: # predicted to rent: 20, # who did rent 3, accuracy 15.0%
K=4: # predicted to rent:  4, # who did rent 0, accuracy 0.0%
K=5: # predicted to rent:  7, # who did rent 1, accuracy 14.3%
```

We see some variability — the numbers for K=4 are very different from the rest.

0.6.2 Comparison to Logistic Regression

As a comparison, we can also fit a logistic regression model to the data. This can also be done with sklearn, though by default it fits something like the ridge regression version of logistic regression, which we introduce in Chapter 6. This can be modified by appropriately setting the argument C below. Its default value is 1 but by setting it to a very large number, the algorithm converges to the same solution as the usual (unregularized) logistic regression estimator discussed above.

Unlike the statsmodels package, sklearn focuses less on inference and more on classification. Hence, the summary methods seen in statsmodels and our simplified version seen with summarize are not generally available for the classifiers in sklearn

```
[66]: logit = LogisticRegression(C=1e10, solver='liblinear')
      logit.fit(X_train, y_train)
      logit_pred = logit.predict_proba(X_test)
      logit_labels = np.where(logit_pred[:,1] > 0.5, 'Yes', 'No')
      confusion_table(logit_labels, y_test)
```

```
[66]: Truth      No  Yes
      Predicted
      No      933  67
      Yes       0   0
```

We used the argument solver='liblinear' above to avoid a warning with the default solver which would indicate that the algorithm does not converge.

If we use 0.5 as the predicted probability cut-of for the classifier, then we have a problem: none of the test observations are predicted to purchase insurance. However, we are not required to use a cut-of of 0.5. If we instead predict a purchase any time the predicted probability of purchase exceeds 0.25, we get much better results: we predict that 29 people will purchase insurance, and we are correct for about 31% of these people. This is almost five times better than random guessing

```
[67]: logit_labels = np.where(logit_pred[:,1]>0.25, 'Yes', 'No')
      confusion_table(logit_labels, y_test)
```

```
[67]: Truth      No  Yes
      Predicted
      No      913   58
      Yes      20    9
```

```
[68]: 9/(20+9)
```

```
[68]: 0.3103448275862069
```

0.7 4.7.7 Linear and Poisson Regression on the Bikeshare Data

Here we fit linear and Poisson regression models to the Bikeshare data, as described in Section 4.6. The response bikers measures the number of bike rentals per hour in Washington, DC in the period 2010–2012.

```
[69]: Bike = load_data('Bikeshare')
```

Let's have a peek at the dimensions and names of the variables in this dataframe.

```
[70]: Bike.shape, Bike.columns
```

```
[70]: ((8645, 15),
      Index(['season', 'mnth', 'day', 'hr', 'holiday', 'weekday', 'workingday',
            'weathersit', 'temp', 'atemp', 'hum', 'windspeed', 'casual',
            'registered', 'bikers'],
            dtype='object'))
```

0.7.1 Linear Regression

We begin by fitting a linear regression model to the data.

```
[71]: X = MS(['mnth', 'hr', 'workingday', 'temp', 'weathersit']).fit_transform(Bike)
      Y = Bike['bikers']
      M_lm = sm.OLS(Y, X).fit()
      summarize(M_lm)
```

```
[71]:
```

	coef	std err	t	P> t
intercept	-68.6317	5.307	-12.932	0.000
mnth[Feb]	6.8452	4.287	1.597	0.110
mnth[March]	16.5514	4.301	3.848	0.000
mnth[April]	41.4249	4.972	8.331	0.000
mnth[May]	72.5571	5.641	12.862	0.000
mnth[June]	67.8187	6.544	10.364	0.000
mnth[July]	45.3245	7.081	6.401	0.000
mnth[Aug]	53.2430	6.640	8.019	0.000
mnth[Sept]	66.6783	5.925	11.254	0.000

mnth[Oct]	75.8343	4.950	15.319	0.000
mnth[Nov]	60.3100	4.610	13.083	0.000
mnth[Dec]	46.4577	4.271	10.878	0.000
hr[1]	-14.5793	5.699	-2.558	0.011
hr[2]	-21.5791	5.733	-3.764	0.000
hr[3]	-31.1408	5.778	-5.389	0.000
hr[4]	-36.9075	5.802	-6.361	0.000
hr[5]	-24.1355	5.737	-4.207	0.000
hr[6]	20.5997	5.704	3.612	0.000
hr[7]	120.0931	5.693	21.095	0.000
hr[8]	223.6619	5.690	39.310	0.000
hr[9]	120.5819	5.693	21.182	0.000
hr[10]	83.8013	5.705	14.689	0.000
hr[11]	105.4234	5.722	18.424	0.000
hr[12]	137.2837	5.740	23.916	0.000
hr[13]	136.0359	5.760	23.617	0.000
hr[14]	126.6361	5.776	21.923	0.000
hr[15]	132.0865	5.780	22.852	0.000
hr[16]	178.5206	5.772	30.927	0.000
hr[17]	296.2670	5.749	51.537	0.000
hr[18]	269.4409	5.736	46.976	0.000
hr[19]	186.2558	5.714	32.596	0.000
hr[20]	125.5492	5.704	22.012	0.000
hr[21]	87.5537	5.693	15.378	0.000
hr[22]	59.1226	5.689	10.392	0.000
hr[23]	26.8376	5.688	4.719	0.000
workingday	1.2696	1.784	0.711	0.477
temp	157.2094	10.261	15.321	0.000
weathersit[cloudy/misty]	-12.8903	1.964	-6.562	0.000
weathersit[heavy rain/snow]	-109.7446	76.667	-1.431	0.152
weathersit[light rain/snow]	-66.4944	2.965	-22.425	0.000

There are 24 levels in `hr` and 40 rows in all, so we have truncated the summary. In `M_lm`, the first levels `hr[0]` and `mnth[Jan]` are treated as the baseline values, and so no coefficient estimates are provided for them: implicitly, their coefficient estimates are zero, and all other levels are measured relative to these baselines. For example, the Feb coefficient of 6.845 signifies that, holding all other variables constant, there are on average about 7 more riders in February than in January. Similarly there are about 16.5 more riders in March than in January.

The results seen in Section 4.6.1 used a slightly different coding of the variables `hr` and `mnth`, as follows:

```
[72]: hr_encode = contrast('hr', 'sum')
      mnth_encode = contrast('mnth', 'sum')
```

Refitting again:


```
[73]: X2 = MS([mnth_encode,hr_encode,'workingday','temp','weathersit']).
      ↪fit_transform(Bike)
M2_lm = sm.OLS(Y, X2).fit()
S2 = summarize(M2_lm)
S2
```

```
[73]:
```

	coef	std err	t	P> t
intercept	73.5974	5.132	14.340	0.000
mnth[Jan]	-46.0871	4.085	-11.281	0.000
mnth[Feb]	-39.2419	3.539	-11.088	0.000
mnth[March]	-29.5357	3.155	-9.361	0.000
mnth[April]	-4.6622	2.741	-1.701	0.089
mnth[May]	26.4700	2.851	9.285	0.000
mnth[June]	21.7317	3.465	6.272	0.000
mnth[July]	-0.7626	3.908	-0.195	0.845
mnth[Aug]	7.1560	3.535	2.024	0.043
mnth[Sept]	20.5912	3.046	6.761	0.000
mnth[Oct]	29.7472	2.700	11.019	0.000
mnth[Nov]	14.2229	2.860	4.972	0.000
hr[0]	-96.1420	3.955	-24.307	0.000
hr[1]	-110.7213	3.966	-27.916	0.000
hr[2]	-117.7212	4.016	-29.310	0.000
hr[3]	-127.2828	4.081	-31.191	0.000
hr[4]	-133.0495	4.117	-32.319	0.000
hr[5]	-120.2775	4.037	-29.794	0.000
hr[6]	-75.5424	3.992	-18.925	0.000
hr[7]	23.9511	3.969	6.035	0.000
hr[8]	127.5199	3.950	32.284	0.000
hr[9]	24.4399	3.936	6.209	0.000
hr[10]	-12.3407	3.936	-3.135	0.002
hr[11]	9.2814	3.945	2.353	0.019
hr[12]	41.1417	3.957	10.397	0.000
hr[13]	39.8939	3.975	10.036	0.000
hr[14]	30.4940	3.991	7.641	0.000
hr[15]	35.9445	3.995	8.998	0.000
hr[16]	82.3786	3.988	20.655	0.000
hr[17]	200.1249	3.964	50.488	0.000
hr[18]	173.2989	3.956	43.806	0.000
hr[19]	90.1138	3.940	22.872	0.000
hr[20]	29.4071	3.936	7.471	0.000
hr[21]	-8.5883	3.933	-2.184	0.029
hr[22]	-37.0194	3.934	-9.409	0.000
workingday	1.2696	1.784	0.711	0.477
temp	157.2094	10.261	15.321	0.000
weathersit[cloudy/misty]	-12.8903	1.964	-6.562	0.000
weathersit[heavy rain/snow]	-109.7446	76.667	-1.431	0.152
weathersit[light rain/snow]	-66.4944	2.965	-22.425	0.000

What is the difference between the two codings? In `M2_lm`, a coefficient estimate is reported for all but level 23 of `hr` and level Dec of `mnth`. Importantly, in `M2_lm`, the (unreported) coefficient estimate for the last level of `mnth` is not zero: instead, it equals the negative of the sum of the coefficient estimates for all of the other levels. Similarly, in `M2_lm`, the coefficient estimate for the last level of `hr` is the negative of the sum of the coefficient estimates for all of the other levels. This means that the coefficients of `hr` and `mnth` in `M2_lm` will always sum to zero, and can be interpreted as the difference from the mean level. For example, the coefficient for January of -46.087 indicates that, holding all other variables constant, there are typically 46 fewer riders in January relative to the yearly average.

It is important to realize that the choice of coding really does not matter, provided that we interpret the model output correctly in light of the coding used. For example, we see that the predictions from the linear model are the same regardless of coding:

```
[74]: np.sum((M_lm.fittedvalues - M2_lm.fittedvalues)**2)
```

```
[74]: 5.2568472069513736e-20
```

The sum of squared differences is zero. We can also see this using the `np.allclose()` function:

```
[75]: np.allclose(M_lm.fittedvalues , M2_lm.fittedvalues)
```

```
[75]: True
```

To reproduce the left-hand side of Figure 4.13 we must first obtain the coefficient estimates associated with `mnth`. The coefficients for January through November can be obtained directly from the `M2_lm` object. The coefficient for December must be explicitly computed as the negative sum of all the other months. We first extract all the coefficients for month from the coefficients of `M2_lm`.

```
[76]: coef_month = S2[S2.index.str.contains('mnth')]['coef']
coef_month
```

```
[76]: mnth[Jan]      -46.0871
      mnth[Feb]     -39.2419
      mnth[March]   -29.5357
      mnth[April]    -4.6622
      mnth[May]      26.4700
      mnth[June]     21.7317
      mnth[July]     -0.7626
      mnth[Aug]       7.1560
      mnth[Sept]     20.5912
      mnth[Oct]      29.7472
      mnth[Nov]      14.2229
      Name: coef, dtype: float64
```

Next, we append Dec as the negative of the sum of all other months.


```
[77]: months = Bike['mnth'].dtype.categories
      coef_month = pd.concat([
      coef_month,
```

```
pd.Series([-coef_month.sum()], index=['mnth[Dec] ']))  
coef_month
```

```
[77]: mnth[Jan]      -46.0871  
      mnth[Feb]     -39.2419  
      mnth[March]   -29.5357  
      mnth[April]    -4.6622  
      mnth[May]      26.4700  
      mnth[June]     21.7317  
      mnth[July]     -0.7626  
      mnth[Aug]       7.1560  
      mnth[Sept]     20.5912  
      mnth[Oct]      29.7472  
      mnth[Nov]      14.2229  
      mnth[Dec]       0.3705  
      dtype: float64
```

Finally, to make the plot neater, we'll just use the first letter of each month, which is the 6th entry of each of the labels in the index.

```
[78]: fig_month, ax_month = subplots(figsize=(8,8))  
      x_month = np.arange(coef_month.shape[0])  
      ax_month.plot(x_month, coef_month, marker='o', ms=10)  
      ax_month.set_xticks(x_month)  
      ax_month.set_xticklabels([l[5] for l in coef_month.index], fontsize  
      =20)  
      ax_month.set_xlabel('Month', fontsize=20)  
      ax_month.set_ylabel('Coefficient', fontsize=20);
```




Reproducing the right-hand plot in Figure 4.13 follows a similar process.

```
[80]: coef_hr = S2[S2.index.str.contains('hr')]['coef']
coef_hr = coef_hr.reindex(['hr[{0}]'.format(h) for h in range(23)])
coef_hr = pd.concat([coef_hr, pd.Series([-coef_hr.sum()], index=['hr[23]'])])
```

We now make the hour plot.

```
[81]: fig_hr, ax_hr = subplots(figsize=(8,8))
x_hr = np.arange(coef_hr.shape[0])
ax_hr.plot(x_hr, coef_hr, marker='o', ms=10)
ax_hr.set_xticks(x_hr[::2])
ax_hr.set_xticklabels(range(24)[::2], fontsize=20)
ax_hr.set_xlabel('Hour', fontsize=20)
```

```
ax_hr.set_ylabel('Coefficient', fontsize=20);
```



0.7.2 Poisson Regression

Now we fit instead a Poisson regression model to the Bikeshare data. Very little changes, except that we now use the function `sm.GLM()` with the Poisson family specified:

```
[82]: M_pois = sm.GLM(Y, X2, family=sm.families.Poisson()).fit()
```

We can plot the coefficients associated with `mnth` and `hr`, in order to reproduce Figure 4.15. We first compute these coefficients as before.


```
[83]: S_pois = summarize(M_pois)
coef_month = S_pois[S_pois.index.str.contains('mnth')]['coef']
```

```
coef_month = pd.concat([coef_month, pd.Series([-coef_month.
    ↪sum()], index=['mnth[Dec]'])])
coef_hr = S_pois[S_pois.index.str.contains('hr')] ['coef']
coef_hr = pd.concat([coef_hr, pd.Series([-coef_hr.sum()], index=['hr[23]'])])
```

The plotting is as before.


```
[84]: fig_pois, (ax_month, ax_hr) = subplots(1, 2, figsize=(16,8))
ax_month.plot(x_month, coef_month, marker='o', ms=10)
ax_month.set_xticks(x_month)
ax_month.set_xticklabels([l[5] for l in coef_month.index], fontsize=20)
ax_month.set_xlabel('Month', fontsize=20)
ax_month.set_ylabel('Coefficient', fontsize=20)
ax_hr.plot(x_hr, coef_hr, marker='o', ms=10)
ax_hr.set_xticklabels(range(24)[::2], fontsize=20)
ax_hr.set_xlabel('Hour', fontsize=20)
ax_hr.set_ylabel('Coefficient', fontsize=20);
```

C:\Users\ankit19.gupta\AppData\Local\Temp\ipykernel_10472\3779510511.py:8:
 UserWarning: set_ticklabels() should only be used with a fixed number of ticks,
 i.e. after set_ticks() or using a FixedLocator.
 ax_hr.set_xticklabels(range(24)[::2], fontsize=20)



We compare the fitted values of the two models. The fitted values are stored in the fittedvalues attribute returned by the fit() method for both the linear regression and the Poisson fts. The linear predictors are stored as the attribute lin_pred.

```
[85]: fig, ax = subplots(figsize=(8, 8))
      ax.scatter(M2_lm.fittedvalues ,M_pois.fittedvalues,s=20)
      ax.set_xlabel('Linear Regression Fit', fontsize=20)
      ax.set_ylabel('Poisson Regression Fit', fontsize=20)
      ax.axline([0,0], c='black', linewidth=3,linestyle='--', slope=1);
```



The predictions from the Poisson regression model are correlated with those from the linear model; however, the former are non-negative. As a result the Poisson regression predictions tend to be larger than those from the linear model for either very low or very high levels of ridership. In this section, we fit Poisson regression models using the `sm.GLM()` function with the argument `family=sm.families.Poisson()`. Earlier in this lab we used the `sm.GLM()` function with `family=sm.families.Binomial()` to perform logistic regression. Other choices for the family argument can be used to fit other types of GLMs. For instance, `family=sm.families.Gamma()` fits a Gamma regression model.

[]: